

Evolving Autoencoders

Christopher Ross Jonathan Brant Zak Roessler


November 1, 2015

Introduction

- Image Processing/Classification
 - ① Each pixel is a separate input
 - ② ANN inputs grows exponentially as resolution increases
 - ③ Method of dimensionality reduction needed
- Autoencoders learns compressed representation
- Traditionally, autoencoder hidden layer topology has been hand-crafted
- Our approach evolves autoencoder structure that can then be optimized
 - ① Connection weights between input/hidden and hidden/output substrate evolved using HyperNEAT
 - ② Backpropagation optimizes evolved autoencoder


Autoencoders

- Autoencoders attempt to extract most salient features from visible layer
- Typical layers include:
 - 1 Input (visible) layer where uncompressed features are received
 - 2 Hidden layer where compressed features are stored
 - 3 Output layer where reconstruction error is calculated
- Most often trained using backpropagation with the intent of minimizing reconstruction error



HyperNEAT

- Indirect encoding substantially reduces search space
 - We're no longer evolving connection weight for each pixel individually
- Substrate configuration hand-designed to capture domain geometry
 - State-space sandwich substrate ideal for visual mapping
- Ability to scale with little-to-no loss of functionality
 - Ideal for image processing at varying resolutions




Evolution of Autoencoders

- Two-phase process:
 - 1 HyperNEAT evolves initial autoencoder connection weights
 - 2 Backpropagation fine tunes autoencoder weights
- Results will be analyzed:
 - 1 Quantitatively - reduction of reconstruction error
 - 2 Qualitatively - how closely does reconstructed image resemble the original
- MNIST handwritten digit dataset will be used for training/validation



Substrate Configuration

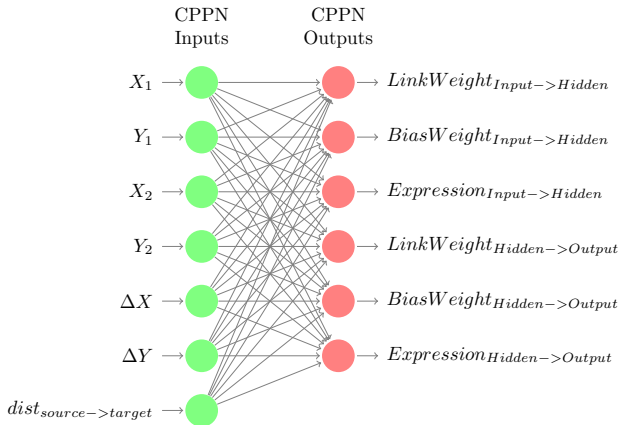
- Three-tiered state-space sandwich substrate
 - 1 Input sheet - all points in 2-dimensional image (size varies based on image resolution)
 - 2 Hidden sheet - compressed feature vector
 - 3 Output sheet - reconstructed image at original resolution



Minimal CPPN

- CPPN inputs are cartesian coordinates between source and target sheet
 - X_1 and Y_1 is point on source sheet
 - X_2 and Y_2 is point on target sheet
 - ΔX and ΔY are difference between X and Y components on source and target sheet
 - $dist_{source \rightarrow target}$ is euclidean distance between source and target point
- CPPN queries the substrate twice:
 - Once for connection weight, bias weight, and expression threshold between input and hidden sheet
 - Once for connection weight, bias weight, and expression threshold between hidden and output sheet

Minimal CPPN





Experiment Configuration

Parameter	Value	Description
Training Sample Proportion	80%	Percentage of sample images used for training
Number of Backpropagation Iterations	100	Number of backpropagation iterations performed during training
Learning Rate	1	Learning Rate (to control backpropagation convergence speed)
Image Resolution Reduction Factor	2	Reduction factor of source image (i.e. a 28x28 image becomes 14x14 with a reduction factor of 2)
MNIST Digit	1	MNIST handwritten digit dataset image used



Current State

- Standard NEAT initially attempted
 - Slow evolution due to massive search space
 - Quantitative reconstruction accuracy rarely exceeded 60%
 - Prone to getting stuck in local optima, yielding to minimal improvements











Current State - HyperNEAT

- HyperNEAT achieved much better reconstruction accuracy
 - Attained over 90% accuracy in a couple of generations
- Matched source images quite closely from a qualitative standpoint



Current State - HyperNEAT

Source	Reconstructed	Set	Index
		Training	30
		Training	60
		Validation	85
		Validation	95